

COMP4030 - Lab 1

Introduction to Jupyter Notebooks and Python for data manipulation

Johann Benerradi
Jeremie Clos

February 2, 2023

Because this is not a programming course, we expect that you have some programming background. Experience with Python is not strictly necessary because we will be using relatively basic features of the language, the most complex of which would be something like list comprehensions. The first lab is intentionally left easy so that students with little experience of Python can take the week to get familiar with the language.

There are many ways of learning Python, depending on your current background and your familiarity with C-type programming languages. There wouldn't be much interest in me copying a Python tutorial into this document, so I will provide useful links instead.

My personal recommendations are the following:

- [Hackingscience.org](https://hackingscience.org) provides a set of 50+ exercises that will get you up to speed on some basic programming in Python.
- The [official Python tutorial](#) is excellent and touches on most if not all features. The essentials sections are 1, 2, 3, 4, 5, 7, 10.5, 10.7
- The [official documentation](#) is extremely instructive and a good webpage to bookmark for future use.
- The [LearnPython](#) tutorial.
- Take [COMP4008-PRG](#) (MSc students only) or buy the book Conceptual Programming with Python by its lecturers.

Jupyter Notebooks, Anaconda, and the CS Virtual Desktop client If you are using your own machine but do not want to use Jupyter Notebooks locally you can make use of the university infrastructure to do the labs, provided you have a solid Internet connection. Follow the steps described on the [University's website](#), for whichever OS your machine is running. Once the setup is complete you can launch the WVD "Computer Science Desktop" from the list of workspaces.

If you want to run Jupyter Notebooks on your own local machine the easiest way is to download and install Anaconda at [Anaconda.com](https://anaconda.com), a scientific Python distribution.

1 Getting started with Jupyter Notebooks

[Jupyter Notebooks](#) are great for teaching and learning Python. They enable to run a Python environment with a graphical interface in your web browser to execute code blocks and write notes in

Markdown language.

You have two main options to run and create Jupyter Notebooks: - I like it simple → use the classic **Jupyter Notebook**: can be found in the CS Virtual Desktop in Anaconda, or can be launched directly from the Windows Start Menu. - I like to customise my developer tools → use **JupyterLab**: can be found in the CS Virtual Desktop in Anaconda

Open the lab's Jupyter Notebook file, **from this point you should be reading the rest of this lab sheet with Jupyter...**

When clicking in the left margin of the content you should notice that the file is composed of many cells, double click on them or hit return when a cell is selected to edit them.

Task 1.1: Get familiar with the interface and find the keyboard shortcuts to: - Create a cell below an existing cell - Create a cell above an existing cell - Delete a cell - Change the type of a cell into a Markdown cell (for note taking) - Change the type of a cell back into a code cell (to execute Python code) - Run a cell (you can run Markdown cells for formatting or code cells for execution)

(Feel free to edit the lab sheet to take notes and write your answers)

[]:

Jupyter Notebooks use the Markdown syntax to format text cells, giving you the opportunity to merge code and documentation into a single executable document. This is an extremely valuable tool for the data and machine learning scientist as it allows you to produce self-documenting computational experiments. You do not have to separately write and code, as you can mix both together. However, Jupyter Notebooks can also be dangerous because they allow you to execute code in a different order than the traditional linear fashion that traditional computer programs follow.

Task 1.2: Get familiar with Markdown and [find the syntax](#) to: - Create headers - Format text (bold/italic) - Insert links - Create block quotes - Create tables

[]:

Most of the libraries useful for this module should already be installed on the CS Virtual Desktop, however there is a way to install a library from a notebook should you need to:

[]:

```
!pip install pandas
```

Note: *Actually the Python interpreter runs within the command prompt, it is a command line interface within a command line interface. In Jupyter we interact with level -1 by default (the Python interpreter) but the ! enables us to interact with the level above, level 0 (so the Windows Command Prompt).*



Important: when submitting work as Jupyter Notebook, always make sure to restart the kernel and re-execute all cells (fast forward button) to be sure that no variables from deleted cells were used.

2 Load data with Pandas

Obviously we first want to...

```
[2]: import pandas as pd
```

Create a DataFrame:

```
[3]: d = {'col1': [1, 2], 'col2': [3, 4]}
df = pd.DataFrame(data=d)
df
```

```
[3]:   col1  col2
0     1     3
1     2     4
```

Create a DataFrame from a file:

```
[4]: with open('countries.csv', 'r') as countries_file:
    headers = countries_file.readline().strip().split(',')
    d = {header: [] for header in headers}
    for line in countries_file:
        l = line.strip().split(',')
        for i, header in enumerate(headers):
            d[header].append(l[i])
    print(d)
df = pd.DataFrame(data=d)
df
```

```
{'COUNTRY': ['China', 'India', 'US', 'Indonesia', 'Brazil', 'Pakistan',
'Nigeria', 'Bangladesh', 'Russia', 'Mexico', 'Japan', 'Germany', 'France', 'UK',
'Italy', 'Argentina', 'Algeria', 'Canada', 'Australia', 'Kazakhstan'], 'POP':
```

```
[1398.72', '1351.16', '329.74', '268.07', '210.32', '205.71', '200.96',
'167.09', '146.79', '126.58', '126.22', '83.02', '67.02', '66.44', '60.36',
'44.94', '43.38', '37.59', '25.47', '18.53'], 'AREA': ['9596.96', '3287.26',
'9833.52', '1910.93', '8515.77', '881.91', '923.77', '147.57', '17098.25',
'1964.38', '377.97', '357.11', '640.68', '242.5', '301.34', '2780.4', '2381.74',
'9984.67', '7692.02', '2724.9'], 'GDP': ['12234.78', '2575.67', '19485.39',
'1015.54', '2055.51', '302.14', '375.77', '245.63', '1530.75', '1158.23',
'4872.42', '3693.2', '2582.49', '2631.23', '1943.84', '637.49', '167.56',
'1647.12', '1408.68', '159.41'], 'CONT': ['Asia', 'Asia', 'N.America', 'Asia',
'S.America', 'Asia', 'Africa', 'Asia', '', 'N.America', 'Asia', 'Europe',
'Europe', 'Europe', 'Europe', 'S.America', 'Africa', 'N.America', 'Oceania',
'Asia']]}
```

```
[4]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
5	Pakistan	205.71	881.91	302.14	Asia
6	Nigeria	200.96	923.77	375.77	Africa
7	Bangladesh	167.09	147.57	245.63	Asia
8	Russia	146.79	17098.25	1530.75	
9	Mexico	126.58	1964.38	1158.23	N.America
10	Japan	126.22	377.97	4872.42	Asia
11	Germany	83.02	357.11	3693.2	Europe
12	France	67.02	640.68	2582.49	Europe
13	UK	66.44	242.5	2631.23	Europe
14	Italy	60.36	301.34	1943.84	Europe
15	Argentina	44.94	2780.4	637.49	S.America
16	Algeria	43.38	2381.74	167.56	Africa
17	Canada	37.59	9984.67	1647.12	N.America
18	Australia	25.47	7692.02	1408.68	Oceania
19	Kazakhstan	18.53	2724.9	159.41	Asia

Actually reinventing the wheel is maybe not that smart...

Task 2.1: There is a Pandas function to create a DataFrame from CSV files, find it in the online documentation and use it to create our DataFrame variable `df_countries` (the table should look the similar to the one above, you should see however that by default missing values are replaced by `NaN`). You may want to have a look at its different arguments for future reference (how to handle headers, different delimiters, ...). Most Python packages have examples in their online documentation which can come in handy.

```
[ ]:
```

```
[6]: df_countries
```

```
[6]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
5	Pakistan	205.71	881.91	302.14	Asia
6	Nigeria	200.96	923.77	375.77	Africa
7	Bangladesh	167.09	147.57	245.63	Asia
8	Russia	146.79	17098.25	1530.75	NaN
9	Mexico	126.58	1964.38	1158.23	N.America
10	Japan	126.22	377.97	4872.42	Asia
11	Germany	83.02	357.11	3693.20	Europe
12	France	67.02	640.68	2582.49	Europe
13	UK	66.44	242.50	2631.23	Europe
14	Italy	60.36	301.34	1943.84	Europe
15	Argentina	44.94	2780.40	637.49	S.America
16	Algeria	43.38	2381.74	167.56	Africa
17	Canada	37.59	9984.67	1647.12	N.America
18	Australia	25.47	7692.02	1408.68	Oceania
19	Kazakhstan	18.53	2724.90	159.41	Asia

3 Select data from tables with Pandas

The `head()` function enables to select the top of a table (this can be handy when working with big tables), by default the 5 first rows:

```
[7]: df_head = df_countries.head()
df_head
```

```
[7]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America

Select columns:

```
[8]: df_head[['COUNTRY', 'GDP']]
```

```
[8]:
```

	COUNTRY	GDP
0	China	12234.78
1	India	2575.67
2	US	19485.39
3	Indonesia	1015.54
4	Brazil	2055.51

Select one column also called a *Series* object (*DataFrame* objects are composed of *Series* objects):

```
[9]: df_head['POP']
```

```
[9]: 0    1398.72
     1    1351.16
     2     329.74
     3     268.07
     4     210.32
     Name: POP, dtype: float64
```

Select rows from indices:

```
[10]: df_countries.iloc[5:10]
```

```
[10]:
```

	COUNTRY	POP	AREA	GDP	CONT
5	Pakistan	205.71	881.91	302.14	Asia
6	Nigeria	200.96	923.77	375.77	Africa
7	Bangladesh	167.09	147.57	245.63	Asia
8	Russia	146.79	17098.25	1530.75	NaN
9	Mexico	126.58	1964.38	1158.23	N.America

Select rows based on a condition:

```
[11]: df_biggest_countries = df_countries[df_countries['AREA'] > 1000] # rows for
     ↪which the area is greater than 1000
df_biggest_countries
```

```
[11]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
8	Russia	146.79	17098.25	1530.75	NaN
9	Mexico	126.58	1964.38	1158.23	N.America
15	Argentina	44.94	2780.40	637.49	S.America
16	Algeria	43.38	2381.74	167.56	Africa
17	Canada	37.59	9984.67	1647.12	N.America
18	Australia	25.47	7692.02	1408.68	Oceania
19	Kazakhstan	18.53	2724.90	159.41	Asia

Select rows and columns based on a condition:

```
[12]: df_countries.loc[df_countries['AREA'] > 1000, ['COUNTRY', 'POP']]
```

```
[12]:
```

	COUNTRY	POP
0	China	1398.72
1	India	1351.16

2	US	329.74
3	Indonesia	268.07
4	Brazil	210.32
8	Russia	146.79
9	Mexico	126.58
15	Argentina	44.94
16	Algeria	43.38
17	Canada	37.59
18	Australia	25.47
19	Kazakhstan	18.53

Select rows and columns **based on their names**:

```
[13]: df_countries_indices = df_countries.rename(index=df_countries['COUNTRY']) # use_
      ↪ countries as row names
      df_countries_indices.loc[['UK', 'France'], ['COUNTRY', 'POP']]
```

```
[13]:      COUNTRY    POP
      UK      UK  66.44
      France France  67.02
```

Select rows and columns **based on indices**:

```
[14]: df_countries.iloc[11:13, 0:2]
```

```
[14]:      COUNTRY    POP
      11 Germany  83.02
      12  France  67.02
```

Task 3.1: Use conditions to store the population of all the European countries of the dataset in a `european_pops` variable.

```
[ ]:
```

```
[16]: list(european_pops)
```

```
[16]: [83.02, 67.02, 66.44, 60.36]
```

4 Arrange tables with Pandas

Task 4.1: Find the Pandas function to sort the `df_countries` DataFrame by descending gross domestic product (GDP), and print only the 10 countries with highest GDP using the `head()` function. Store this sorted table in `df_sort_1`.

```
[ ]:
```

```
[18]: df_sort_1
```

```
[18]:
```

	COUNTRY	POP	AREA	GDP	CONT
2	US	329.74	9833.52	19485.39	N.America
0	China	1398.72	9596.96	12234.78	Asia
10	Japan	126.22	377.97	4872.42	Asia
11	Germany	83.02	357.11	3693.20	Europe
13	UK	66.44	242.50	2631.23	Europe
12	France	67.02	640.68	2582.49	Europe
1	India	1351.16	3287.26	2575.67	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
14	Italy	60.36	301.34	1943.84	Europe
17	Canada	37.59	9984.67	1647.12	N.America

Task 4.2: Use the same Pandas sorting function to sort the `df_countries` DataFrame by continent's alphabetical order and descending area. Store this sorted table in `df_sort_2`. Take a note of how *NaN* values are handled, you can change this behaviour with the `na_position` argument.

```
[ ]:
```

```
[20]: df_sort_2
```

```
[20]:
```

	COUNTRY	POP	AREA	GDP	CONT
16	Algeria	43.38	2381.74	167.56	Africa
6	Nigeria	200.96	923.77	375.77	Africa
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
19	Kazakhstan	18.53	2724.90	159.41	Asia
3	Indonesia	268.07	1910.93	1015.54	Asia
5	Pakistan	205.71	881.91	302.14	Asia
10	Japan	126.22	377.97	4872.42	Asia
7	Bangladesh	167.09	147.57	245.63	Asia
12	France	67.02	640.68	2582.49	Europe
11	Germany	83.02	357.11	3693.20	Europe
14	Italy	60.36	301.34	1943.84	Europe
13	UK	66.44	242.50	2631.23	Europe
17	Canada	37.59	9984.67	1647.12	N.America
2	US	329.74	9833.52	19485.39	N.America
9	Mexico	126.58	1964.38	1158.23	N.America
18	Australia	25.47	7692.02	1408.68	Oceania
4	Brazil	210.32	8515.77	2055.51	S.America
15	Argentina	44.94	2780.40	637.49	S.America
8	Russia	146.79	17098.25	1530.75	NaN

Assign new values using indices:

```
[21]: df_countries_post_2020 = df_countries.copy()
df_countries_post_2020.loc['UK', 'CONT'] = 'Brexit'
df_countries_post_2020
```



```
[21]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
5	Pakistan	205.71	881.91	302.14	Asia
6	Nigeria	200.96	923.77	375.77	Africa
7	Bangladesh	167.09	147.57	245.63	Asia
8	Russia	146.79	17098.25	1530.75	NaN
9	Mexico	126.58	1964.38	1158.23	N.America
10	Japan	126.22	377.97	4872.42	Asia
11	Germany	83.02	357.11	3693.20	Europe
12	France	67.02	640.68	2582.49	Europe
13	UK	66.44	242.50	2631.23	Europe
14	Italy	60.36	301.34	1943.84	Europe
15	Argentina	44.94	2780.40	637.49	S.America
16	Algeria	43.38	2381.74	167.56	Africa
17	Canada	37.59	9984.67	1647.12	N.America
18	Australia	25.47	7692.02	1408.68	Oceania
19	Kazakhstan	18.53	2724.90	159.41	Asia
UK	NaN	NaN	NaN	NaN	Brexit

Create a new column from another column:

```
[22]: df_countries['AREA_SQUARE_KM'] = df_countries['AREA'] * 1000
df_countries
```

```
[22]:
```

	COUNTRY	POP	AREA	GDP	CONT	AREA_SQUARE_KM
0	China	1398.72	9596.96	12234.78	Asia	9596960.0
1	India	1351.16	3287.26	2575.67	Asia	3287260.0
2	US	329.74	9833.52	19485.39	N.America	9833520.0
3	Indonesia	268.07	1910.93	1015.54	Asia	1910930.0
4	Brazil	210.32	8515.77	2055.51	S.America	8515770.0
5	Pakistan	205.71	881.91	302.14	Asia	881910.0
6	Nigeria	200.96	923.77	375.77	Africa	923770.0
7	Bangladesh	167.09	147.57	245.63	Asia	147570.0
8	Russia	146.79	17098.25	1530.75	NaN	17098250.0
9	Mexico	126.58	1964.38	1158.23	N.America	1964380.0
10	Japan	126.22	377.97	4872.42	Asia	377970.0
11	Germany	83.02	357.11	3693.20	Europe	357110.0
12	France	67.02	640.68	2582.49	Europe	640680.0
13	UK	66.44	242.50	2631.23	Europe	242500.0
14	Italy	60.36	301.34	1943.84	Europe	301340.0
15	Argentina	44.94	2780.40	637.49	S.America	2780400.0
16	Algeria	43.38	2381.74	167.56	Africa	2381740.0
17	Canada	37.59	9984.67	1647.12	N.America	9984670.0
18	Australia	25.47	7692.02	1408.68	Oceania	7692020.0

19	Kazakhstan	18.53	2724.90	159.41	Asia	2724900.0
----	------------	-------	---------	--------	------	-----------

Delete a column:

```
[23]: df_countries = df_countries.drop(columns=['AREA'])
df_countries
```

```
[23]:
```

	COUNTRY	POP	GDP	CONT	AREA_SQUARE_KM
0	China	1398.72	12234.78	Asia	9596960.0
1	India	1351.16	2575.67	Asia	3287260.0
2	US	329.74	19485.39	N.America	9833520.0
3	Indonesia	268.07	1015.54	Asia	1910930.0
4	Brazil	210.32	2055.51	S.America	8515770.0
5	Pakistan	205.71	302.14	Asia	881910.0
6	Nigeria	200.96	375.77	Africa	923770.0
7	Bangladesh	167.09	245.63	Asia	147570.0
8	Russia	146.79	1530.75	NaN	17098250.0
9	Mexico	126.58	1158.23	N.America	1964380.0
10	Japan	126.22	4872.42	Asia	377970.0
11	Germany	83.02	3693.20	Europe	357110.0
12	France	67.02	2582.49	Europe	640680.0
13	UK	66.44	2631.23	Europe	242500.0
14	Italy	60.36	1943.84	Europe	301340.0
15	Argentina	44.94	637.49	S.America	2780400.0
16	Algeria	43.38	167.56	Africa	2381740.0
17	Canada	37.59	1647.12	N.America	9984670.0
18	Australia	25.47	1408.68	Oceania	7692020.0
19	Kazakhstan	18.53	159.41	Asia	2724900.0

Task 4.3: Assigning new values can be useful in cases where we want to discretise quantitative data. Create a new column SIZE where the value is Large if the area of the country is greater than 1,000,000 square km and Small if less.

```
[ ]:
```

```
[25]: df_countries
```

```
[25]:
```

	COUNTRY	POP	GDP	CONT	AREA_SQUARE_KM	SIZE
0	China	1398.72	12234.78	Asia	9596960.0	Large
1	India	1351.16	2575.67	Asia	3287260.0	Large
2	US	329.74	19485.39	N.America	9833520.0	Large
3	Indonesia	268.07	1015.54	Asia	1910930.0	Large
4	Brazil	210.32	2055.51	S.America	8515770.0	Large
5	Pakistan	205.71	302.14	Asia	881910.0	Small
6	Nigeria	200.96	375.77	Africa	923770.0	Small
7	Bangladesh	167.09	245.63	Asia	147570.0	Small
8	Russia	146.79	1530.75	NaN	17098250.0	Large
9	Mexico	126.58	1158.23	N.America	1964380.0	Large

10	Japan	126.22	4872.42	Asia	377970.0	Small
11	Germany	83.02	3693.20	Europe	357110.0	Small
12	France	67.02	2582.49	Europe	640680.0	Small
13	UK	66.44	2631.23	Europe	242500.0	Small
14	Italy	60.36	1943.84	Europe	301340.0	Small
15	Argentina	44.94	637.49	S.America	2780400.0	Large
16	Algeria	43.38	167.56	Africa	2381740.0	Large
17	Canada	37.59	1647.12	N.America	9984670.0	Large
18	Australia	25.47	1408.68	Oceania	7692020.0	Large
19	Kazakhstan	18.53	159.41	Asia	2724900.0	Large

Write table to a file:

```
[26]: df_countries.to_csv('countries_new.csv', index=False)
```

5 Combine tables with Pandas

Read the file in two halves:

```
[27]: df_1 = pd.read_csv('countries.csv', nrows=15)
      df_2 = pd.read_csv('countries.csv', skiprows=range(1,16))
```

```
[28]: print(df_1.shape)
      df_1
```

(15, 5)

```
[28]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
5	Pakistan	205.71	881.91	302.14	Asia
6	Nigeria	200.96	923.77	375.77	Africa
7	Bangladesh	167.09	147.57	245.63	Asia
8	Russia	146.79	17098.25	1530.75	NaN
9	Mexico	126.58	1964.38	1158.23	N.America
10	Japan	126.22	377.97	4872.42	Asia
11	Germany	83.02	357.11	3693.20	Europe
12	France	67.02	640.68	2582.49	Europe
13	UK	66.44	242.50	2631.23	Europe
14	Italy	60.36	301.34	1943.84	Europe

```
[29]: print(df_2.shape)
      df_2
```

(5, 5)

```
[29]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	Argentina	44.94	2780.40	637.49	S.America
1	Algeria	43.38	2381.74	167.56	Africa
2	Canada	37.59	9984.67	1647.12	N.America
3	Australia	25.47	7692.02	1408.68	Oceania
4	Kazakhstan	18.53	2724.90	159.41	Asia

Combine tables vertically (add rows):

```
[30]: df = pd.concat([df_1, df_2], axis=0).reset_index(drop=True)
df
```

```
[30]:
```

	COUNTRY	POP	AREA	GDP	CONT
0	China	1398.72	9596.96	12234.78	Asia
1	India	1351.16	3287.26	2575.67	Asia
2	US	329.74	9833.52	19485.39	N.America
3	Indonesia	268.07	1910.93	1015.54	Asia
4	Brazil	210.32	8515.77	2055.51	S.America
5	Pakistan	205.71	881.91	302.14	Asia
6	Nigeria	200.96	923.77	375.77	Africa
7	Bangladesh	167.09	147.57	245.63	Asia
8	Russia	146.79	17098.25	1530.75	NaN
9	Mexico	126.58	1964.38	1158.23	N.America
10	Japan	126.22	377.97	4872.42	Asia
11	Germany	83.02	357.11	3693.20	Europe
12	France	67.02	640.68	2582.49	Europe
13	UK	66.44	242.50	2631.23	Europe
14	Italy	60.36	301.34	1943.84	Europe
15	Argentina	44.94	2780.40	637.49	S.America
16	Algeria	43.38	2381.74	167.56	Africa
17	Canada	37.59	9984.67	1647.12	N.America
18	Australia	25.47	7692.02	1408.68	Oceania
19	Kazakhstan	18.53	2724.90	159.41	Asia

Combine tables horizontally (add columns):

```
[31]: df_A = df[['COUNTRY', 'GDP']].sort_values(by=['GDP']).iloc[2:]
df_A
```

```
[31]:
```

	COUNTRY	GDP
7	Bangladesh	245.63
5	Pakistan	302.14
6	Nigeria	375.77
15	Argentina	637.49
3	Indonesia	1015.54
9	Mexico	1158.23
18	Australia	1408.68
8	Russia	1530.75

17	Canada	1647.12
14	Italy	1943.84
4	Brazil	2055.51
1	India	2575.67
12	France	2582.49
13	UK	2631.23
11	Germany	3693.20
10	Japan	4872.42
0	China	12234.78
2	US	19485.39

```
[32]: df_B = df_2[['COUNTRY', 'AREA', 'POP']]
      df_B
```

```
[32]:
```

	COUNTRY	AREA	POP
0	Argentina	2780.40	44.94
1	Algeria	2381.74	43.38
2	Canada	9984.67	37.59
3	Australia	7692.02	25.47
4	Kazakhstan	2724.90	18.53

```
[33]: pd.merge(df_A, df_B, on='COUNTRY', how='inner')
```

```
[33]:
```

	COUNTRY	GDP	AREA	POP
0	Argentina	637.49	2780.40	44.94
1	Australia	1408.68	7692.02	25.47
2	Canada	1647.12	9984.67	37.59

Task 5.1: have a look at documentation and try to understand the behaviour of the different possible values of the `how` argument of that `merge()` function.

```
[ ]:
```

6 Get basic statistics with Pandas

Summary statistics on the table:

```
[34]: df.describe()
```

```
[34]:
```

	POP	AREA	GDP
count	20.000000	20.000000	20.000000
mean	248.905500	4082.182500	3036.142500
std	394.546143	4706.507539	4706.00783
min	18.530000	147.570000	159.41000
25%	56.505000	575.002500	572.06000
50%	126.400000	2173.060000	1588.93500
75%	206.862500	7897.957500	2594.67500
max	1398.720000	17098.250000	19485.39000

```
[35]: df.median(numeric_only=True)
```

```
[35]: POP      126.400  
      AREA    2173.060  
      GDP     1588.935  
      dtype: float64
```

```
[36]: df['POP'].skew()
```

```
[36]: 2.6479537986199797
```

```
[37]: df['GDP'].kurtosis()
```

```
[37]: 8.356678932041959
```

Task 6.1: do some research on what all those statistics represent if you don't know about them.

```
[ ]:
```

7 Final tasks

Final task 1: Using the dataset of this lab, present in a new column of the table the density of population of every country in people per square meter.

```
[ ]:
```

Final task 2: What continent has the countries with the largest density of population on average?

```
[ ]:
```

Final task 3: Why is it however not scientifically valid to answer this previous research question with our dataset?

```
[ ]:
```

8 Open tasks

Tasks are more time-consuming and open-ended than exercises and allow you to go much deeper in the mastery of the content. We do not expect you to finish all of them in the lab, but they can be good toy problems to practice with.

Open task 1: Load the [Penguin dataset](#) and try to perform similar data manipulations as what we have seen on it.

```
[38]: # The Penguin dataset can be loaded with seaborn (data visualisation library for  
      ↪ Python)  
      # We will use this library more in depth in later labs  
  
import seaborn as sns
```

```
df_penguins = sns.load_dataset('penguins')
df_penguins
```

```
[38]:
```

	species	island	bill_length_mm	bill_depth_mm	flipper_length_mm	\
0	Adelie	Torgersen	39.1	18.7	181.0	
1	Adelie	Torgersen	39.5	17.4	186.0	
2	Adelie	Torgersen	40.3	18.0	195.0	
3	Adelie	Torgersen	NaN	NaN	NaN	
4	Adelie	Torgersen	36.7	19.3	193.0	
...	
339	Gentoo	Biscoe	NaN	NaN	NaN	
340	Gentoo	Biscoe	46.8	14.3	215.0	
341	Gentoo	Biscoe	50.4	15.7	222.0	
342	Gentoo	Biscoe	45.2	14.8	212.0	
343	Gentoo	Biscoe	49.9	16.1	213.0	

	body_mass_g	sex
0	3750.0	Male
1	3800.0	Female
2	3250.0	Female
3	NaN	NaN
4	3450.0	Female
...
339	NaN	NaN
340	4850.0	Female
341	5750.0	Male
342	5200.0	Female
343	5400.0	Male

[344 rows x 7 columns]

```
[ ]:
```

Open task 2: Get familiar with the `pivot()` and `melt()` functions from Pandas.

```
[ ]:
```